

Data collection guide

Free and public. No login, no API. Budget about an hour.

Which report to use

Tradestat offers eleven report types and the choice matters enormously:

Report	Gives you	Downloads needed
Commodity x Country-wise	One HS code, one country	11 x 200+ x 8 — unusable
Commodity wise all Countries	One HS code, all countries , one year	11 x 8 = 88
Chapter-wise all commodities	Chapter totals only, no destinations	—

Use **Commodity wise all Countries**. The other country-level report makes you pick a single destination per query, which would mean thousands of downloads.

1. Annual data (EIDB)

https://tradestat.commerce.gov.in/eidb/commodity_wise_all_countries_export

Coverage: 2018-19 to 2025-26 (eight years in the dropdown).

The form has exactly three inputs:

HS Code — type the 4-digit heading, or use the *Search HSCode* box

Year — one financial year per query

Values in — choose **US \$ Million**, not Rs Crore

Hit **Submit**, then export the result table.

Naming your downloads — this part is not optional

The report returns one HS code for one year, so **neither the HS code nor the year appears as a column in the file**. They exist only in the filename. The parser reads them from there.

Save every file as `<hscode>_<year>.csv`:



5205_2018-19.csv
5205_2019-20.csv
5402_2018-19.csv

Get this wrong and the pipeline will reject the file with an error telling you to rename it.

Which headings

Do not use chapter totals. Chapters 52, 54 and 55 include raw fibre, fabric and waste — chapter figures would be far larger than actual yarn exports and the whole analysis would be wrong.

Start with these three. They carry most of the value, and 3 x 8 = 24 downloads is one sitting:

Heading	What it is
5205	Cotton yarn, >=85% cotton, not retail
5402	Synthetic filament yarn, not retail
5509	Yarn of synthetic staple fibres, not retail

Add later for the complete picture: 5206, 5207 (cotton), 5401, 5403, 5406 (filament), 5508, 5510, 5511 (staple).

Run the pipeline after the first three. If the output looks right, carry on; if not, you have wasted 24 downloads rather than 88.

2. Monthly data (MEIDB) — for seasonality

https://tradedat.commerce.gov.in/meidb/commoditywise_export

Coverage: Jan 2018 to Mar 2026. Note this is **calendar** months, whereas EIDB uses April-March financial years. Keep the two sets separate; the pipeline already routes them to different analyses.

Same approach, same naming convention. Annual data cannot show seasonality, and seasonality is one of the more commercially interesting findings here — yarn shipments are not flat across the year.

3. Competitor benchmark (UN Comtrade)

Tradedat only knows what India ships. To answer "is India gaining or losing share against China and Vietnam in the same headings", you need mirror data:

☐

```
py src\fetch_comtrade.py
```

The free public preview endpoint needs no authorisation. It is rate-limited, so the script pauses between calls. If you hit limits, register a free key at <https://comtradedeveloper.un.org/> (auto-approved) and set it first:

☐

```
set COMTRADE_KEY=your_key_here
```

4. Optional: port-wise data

Tradestat publishes port-wise trade under the FTSPCC section. Pulling it makes the trade-lane map genuinely original — very few analyses look at which Indian ports serve which destinations, and it is something you can speak to from experience.

Data quality notes

Things that will bite you:

Country names are DGCI&S abbreviations, not standard names: `BANGLADESH PR`, `CHINA P RP`, `EGYPT A RP`, `KOREA RP`, `U ARAB EMTS`, `U K`, `U S A`, `VIETNAM SOC REP`, `SRI LANKA DSR`, `PAKISTAN IR`, `SAUDI ARAB`, `NETHERLAND`. The common ones are mapped in `visualise.NAME_FIXES`; add any others you meet.

`UNSPECIFIED` and `Trade to Unspecified Countries` **are not destinations**. They are residual buckets and are dropped automatically.

Financial years run April-March. `2023-24` = Apr 2023 to Mar 2024. Not directly comparable to Comtrade calendar years.

Nil is written as `-`, not 0 or blank. Handled.

Total rows sit inside the table. Dropped automatically, but sanity- check your row counts.

HS codes were re-allocated from April 2026 — the portal flags this. If a heading's value jumps oddly in the final year, that is likely why.

The latest year is provisional and gets revised. Say so in your write-up; noting it shows you understand the source.

Checklist

- ☐ 24 files (3 headings x 8 years) named `<hscode>_<year>.csv`
- ☐ All in US \$ Million, not Rs Crore

- ☐ Saved to `data\raw\`
- ☐ `SYNTHETIC_*` files deleted
- ☐ `py_src\run_all.py` runs without errors